

Teads

**Moving from Pull to Push
and dividing by 1000 the network traffic**

Sunny Tech 2023

Agenda

Context

Architecture

Implementation

Feedback



Context Business

Advertising



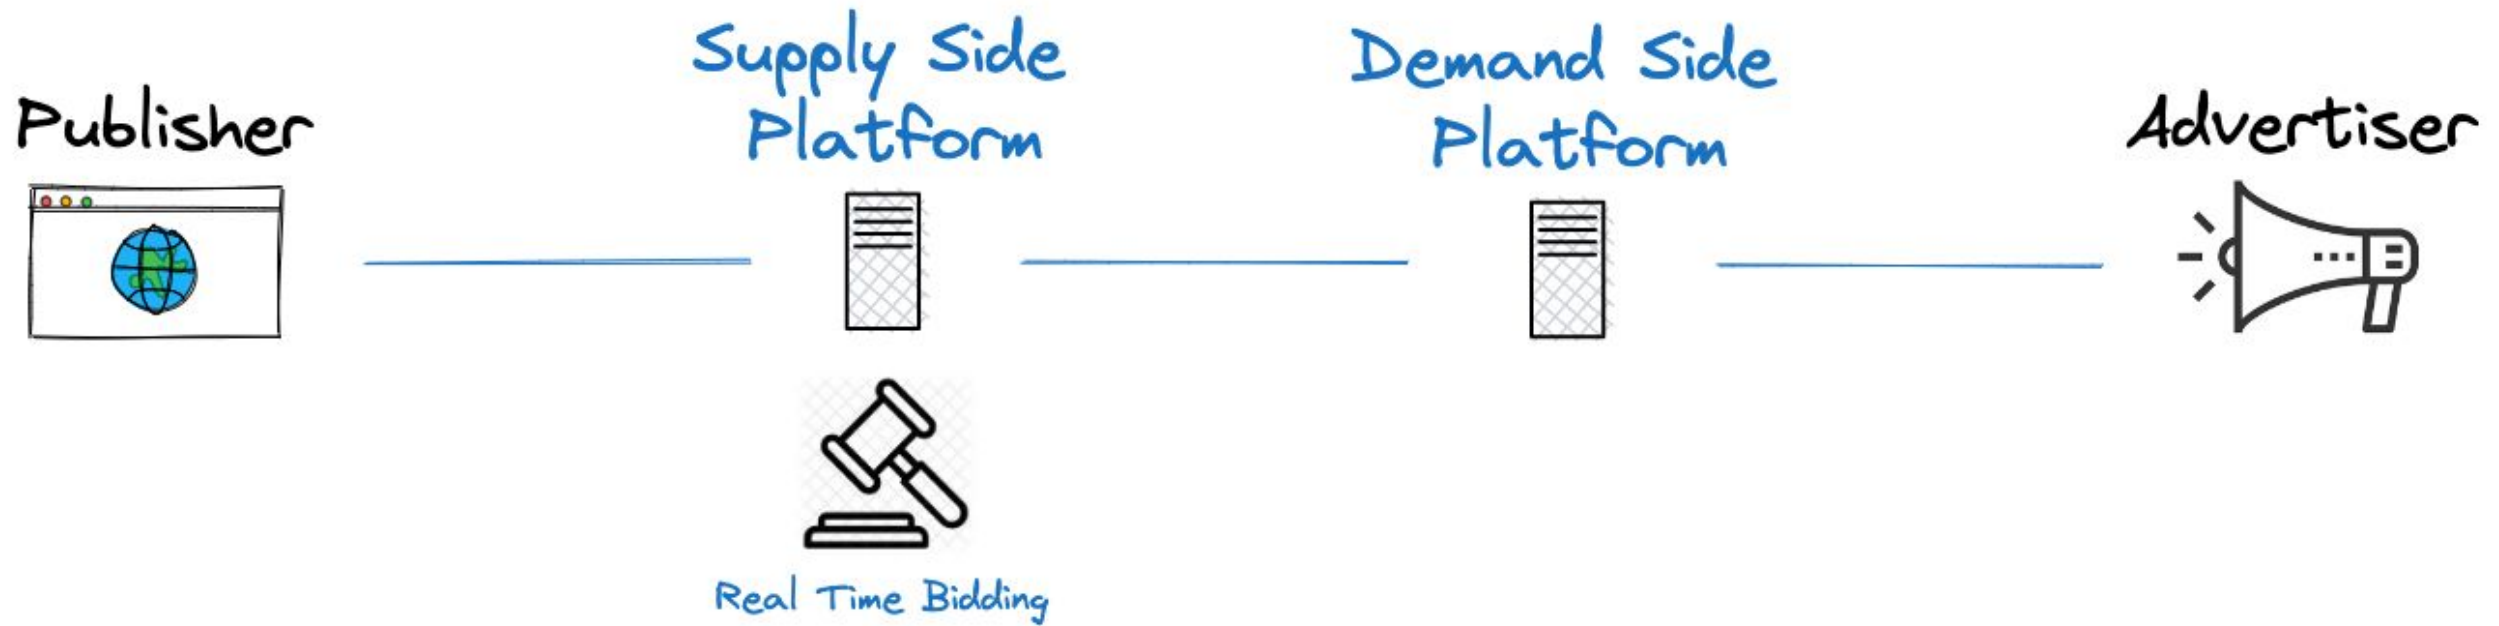
Digital Advertising at the beginnings



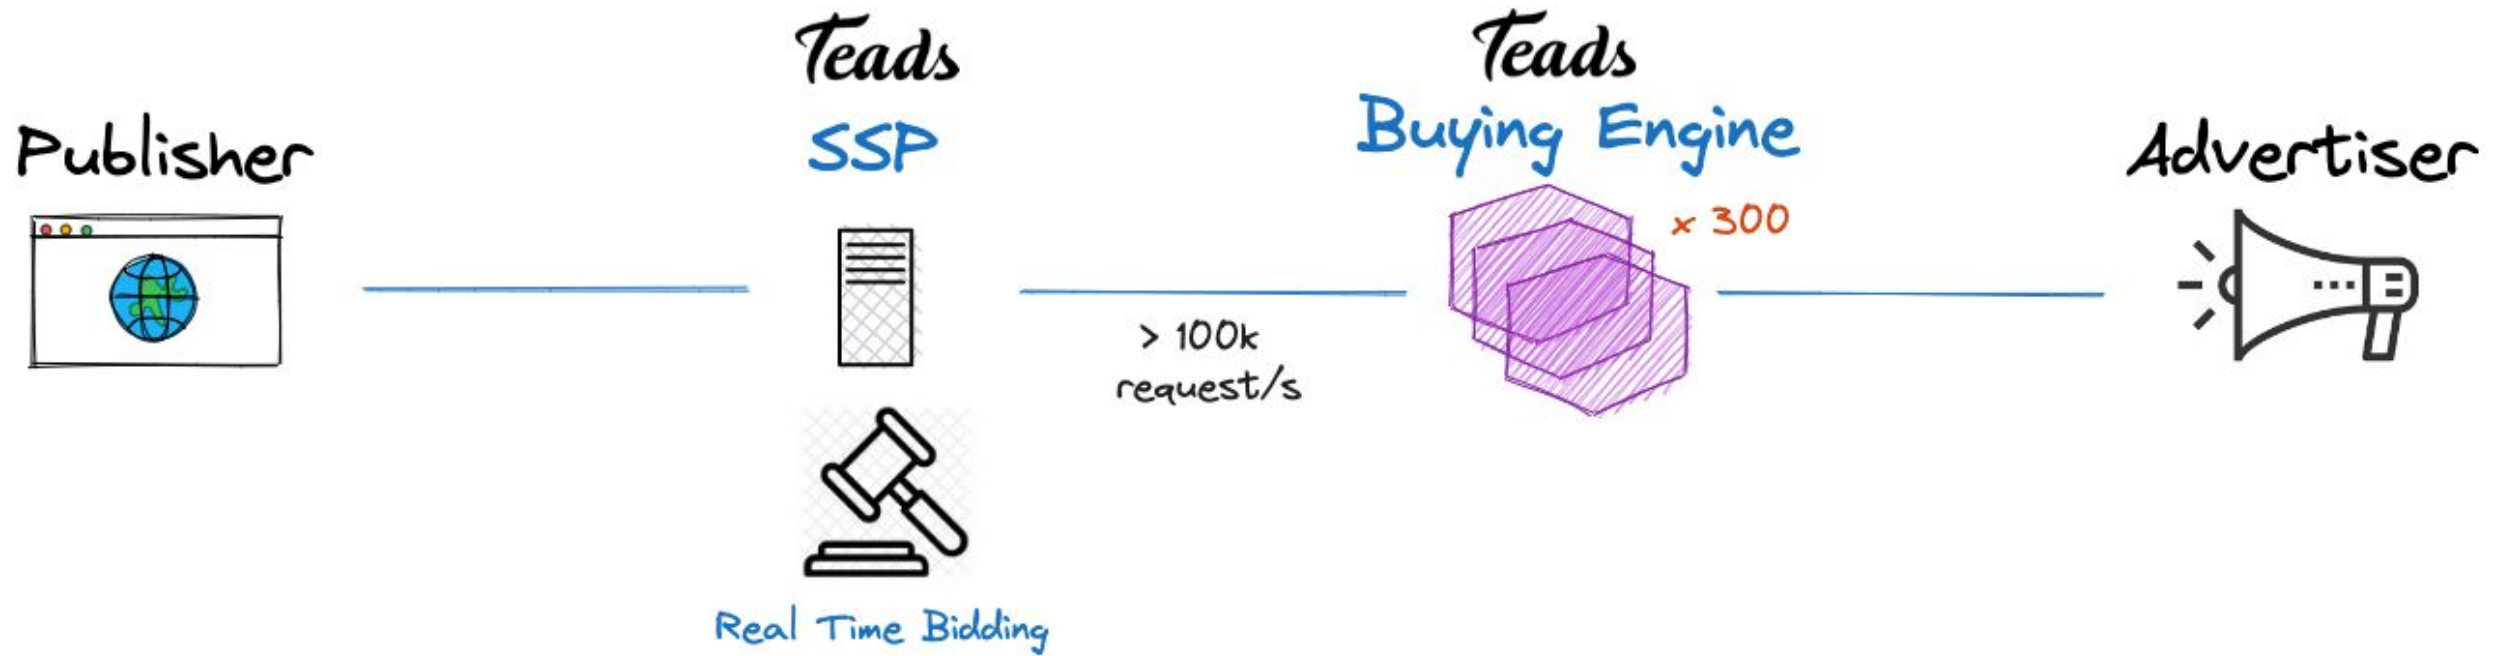
Digital Advertising in the 2000s



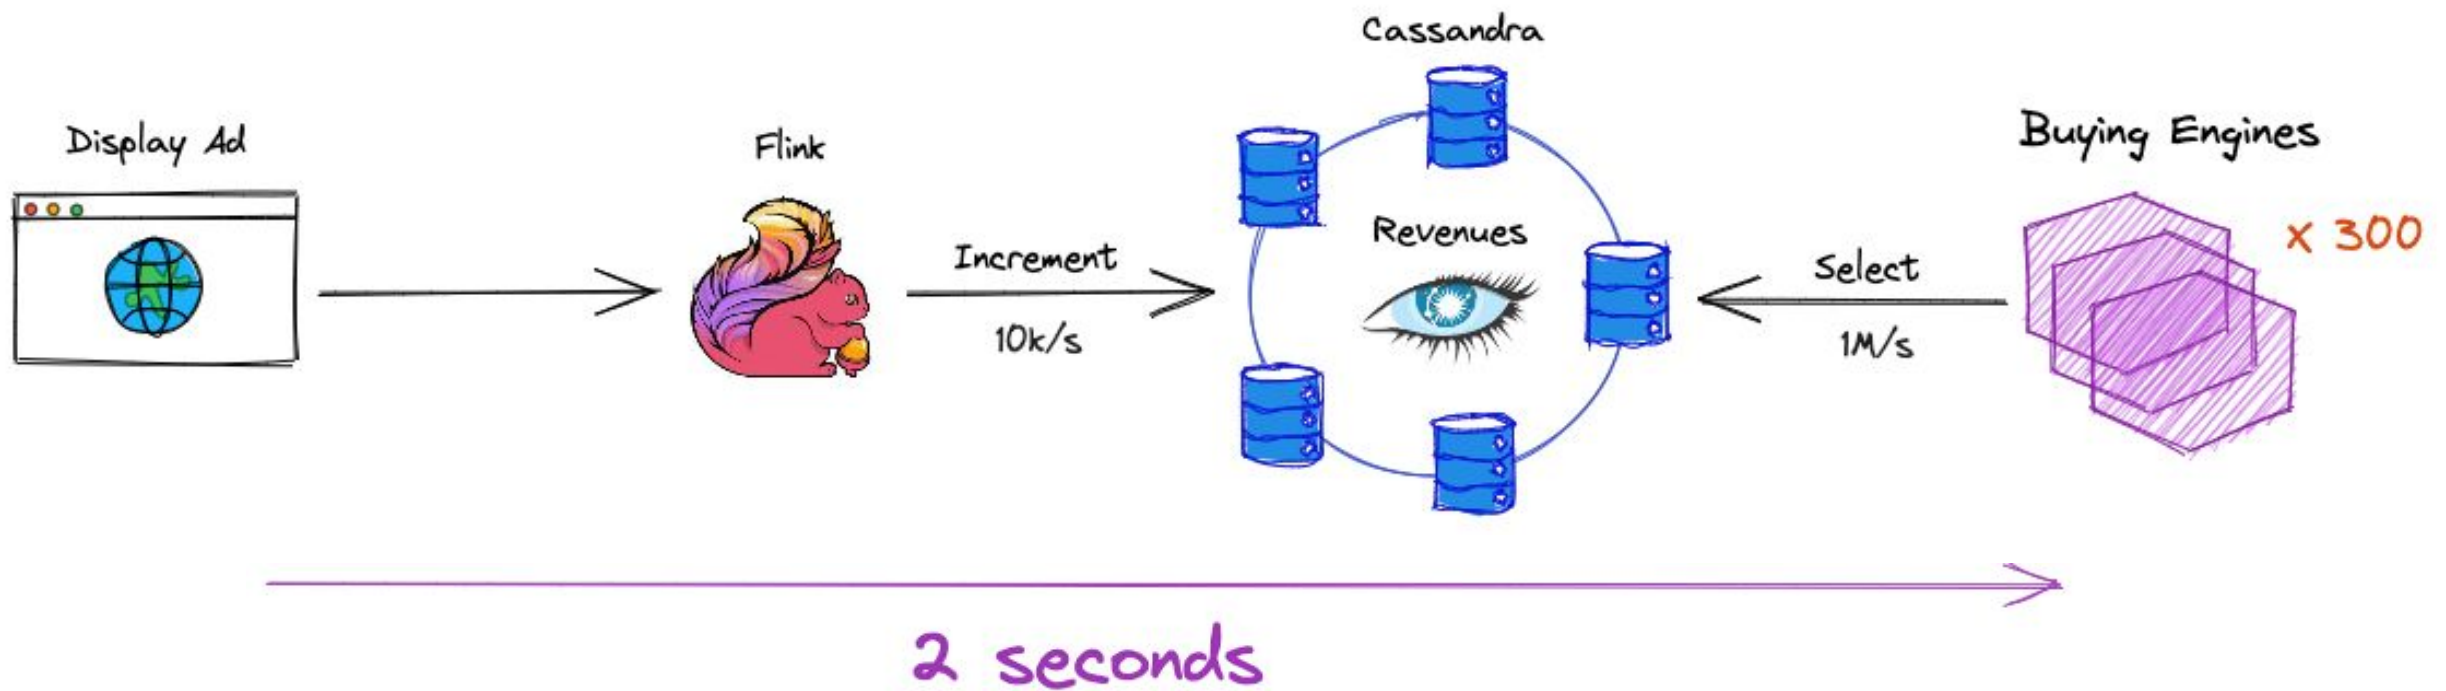
Digital Advertising in the 2010s



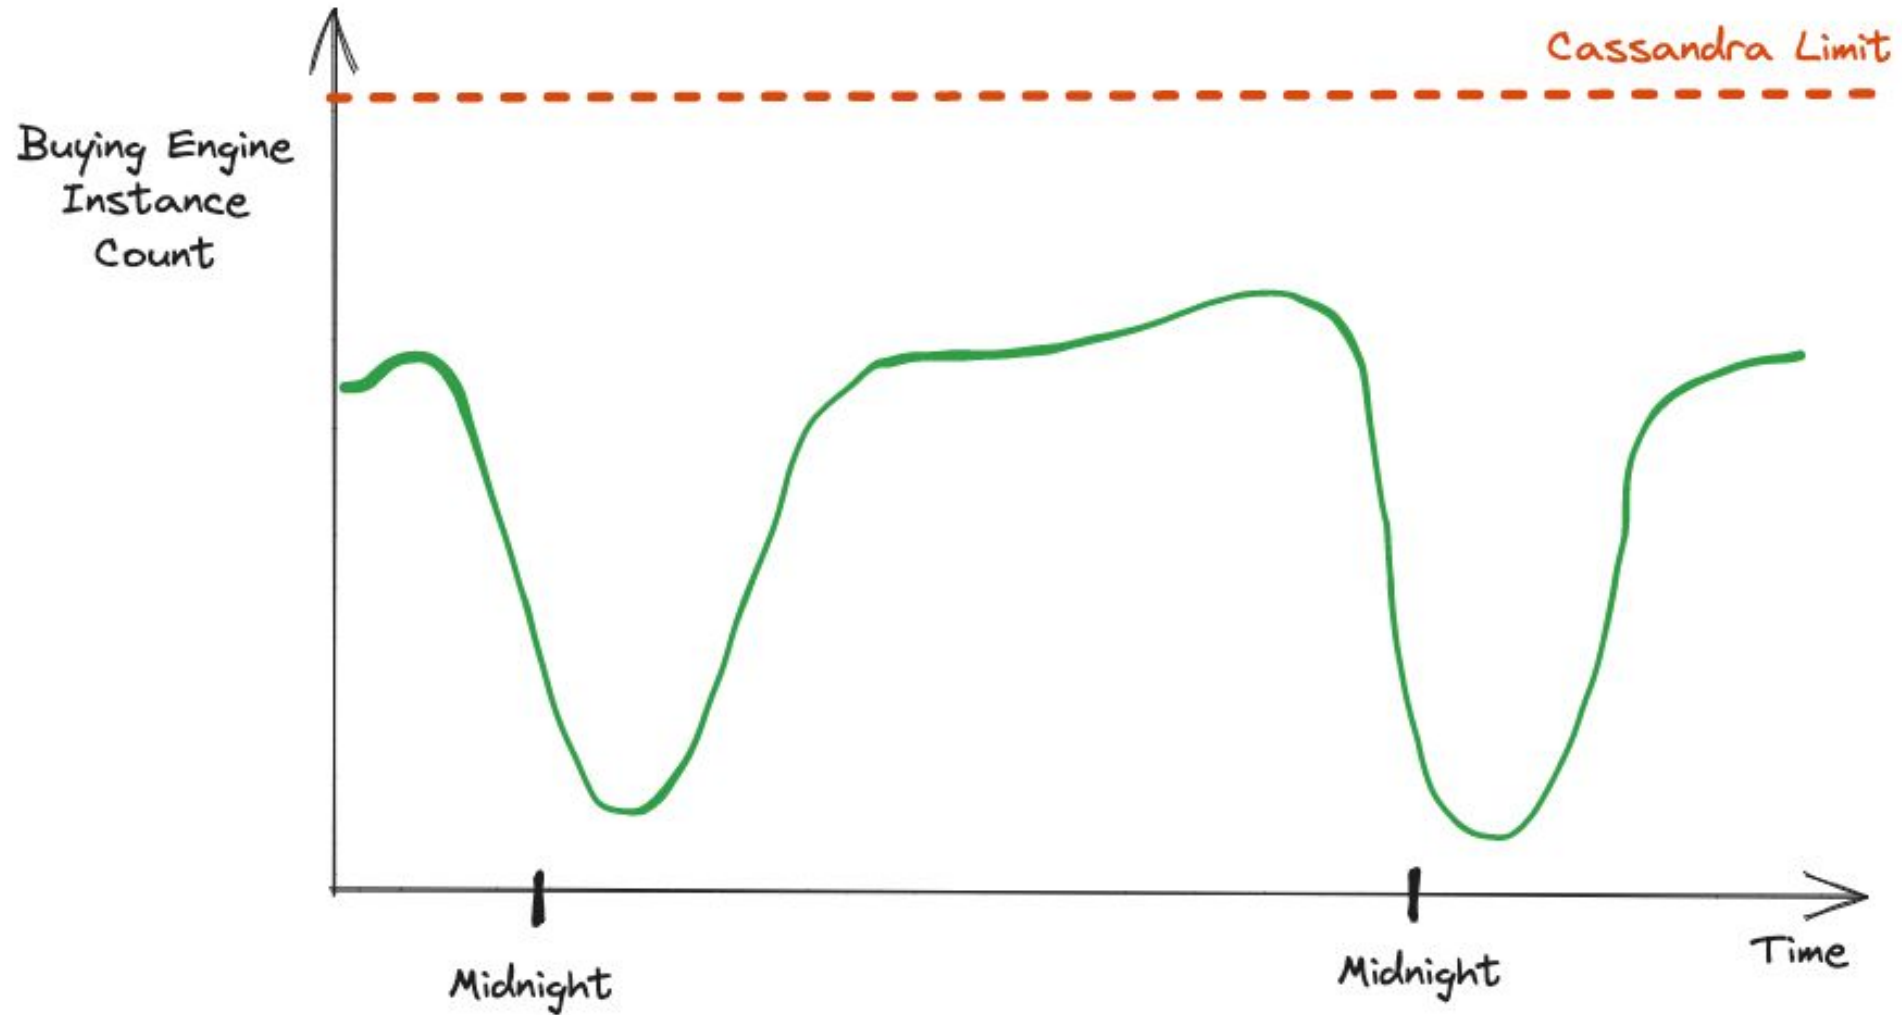
The Buying Engine



Ad Budget Control



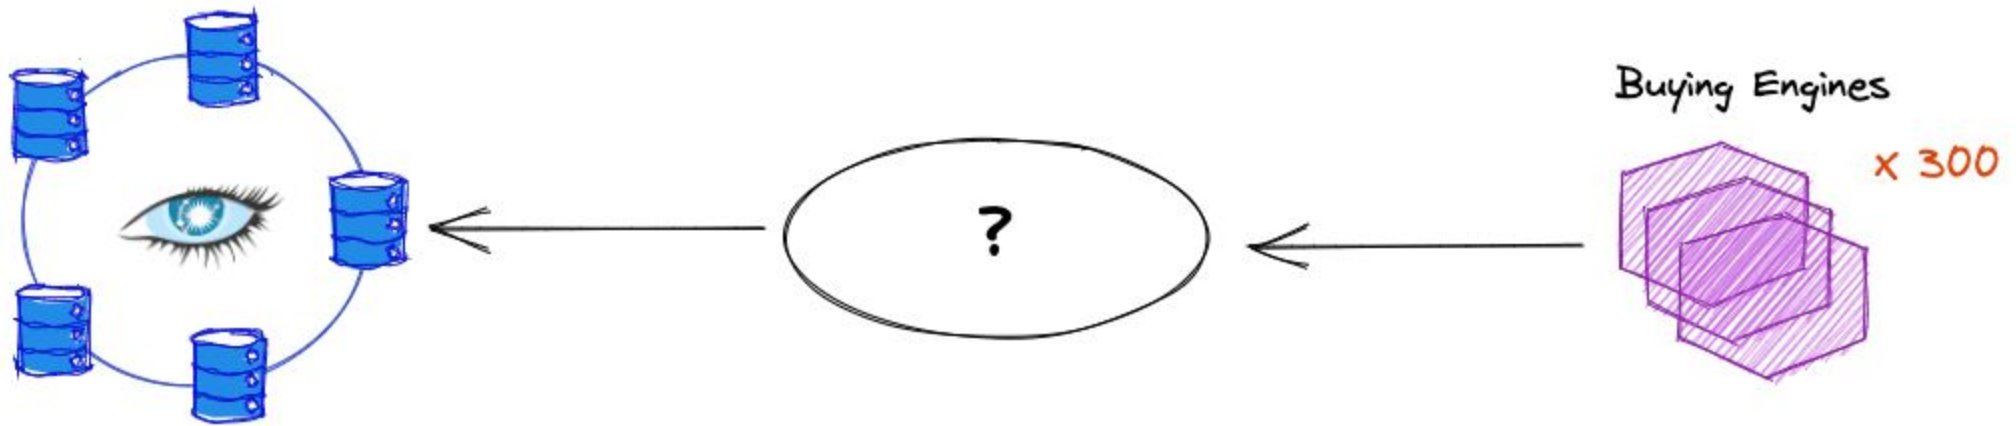
The Scaling Limitation



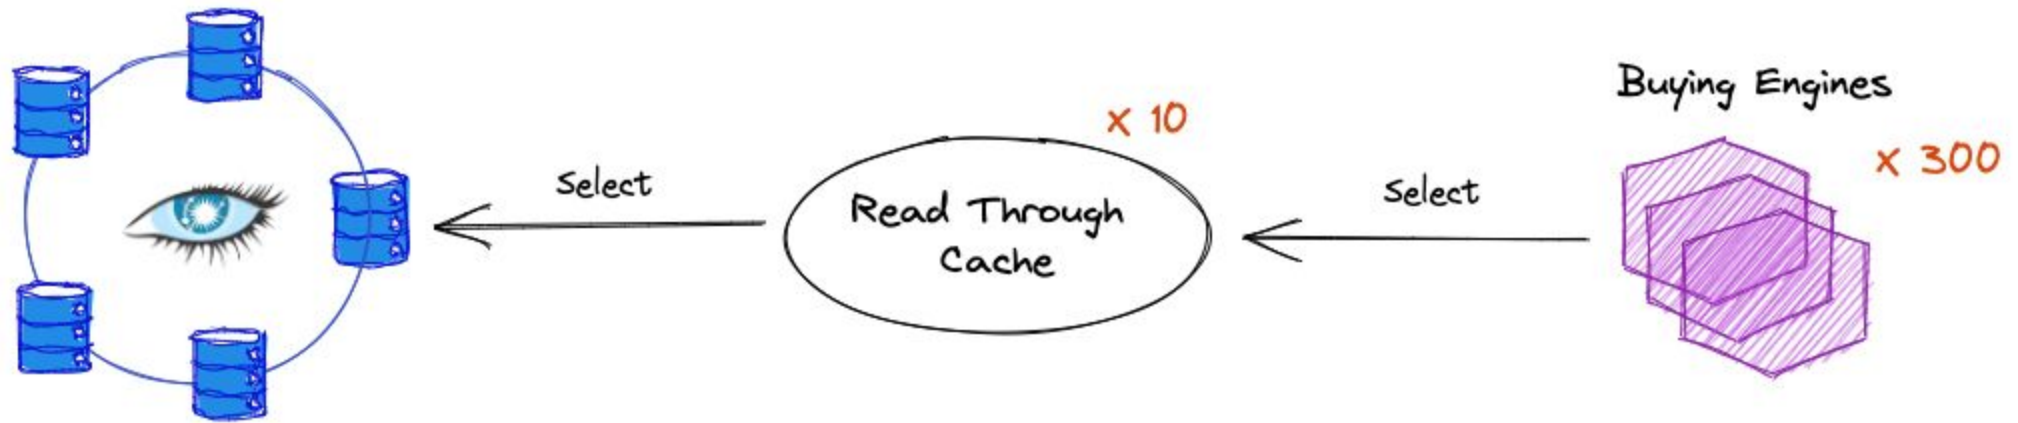


Architecture Design

Decorrelate Cassandra & Buying Engines



A Read Through Cache ?



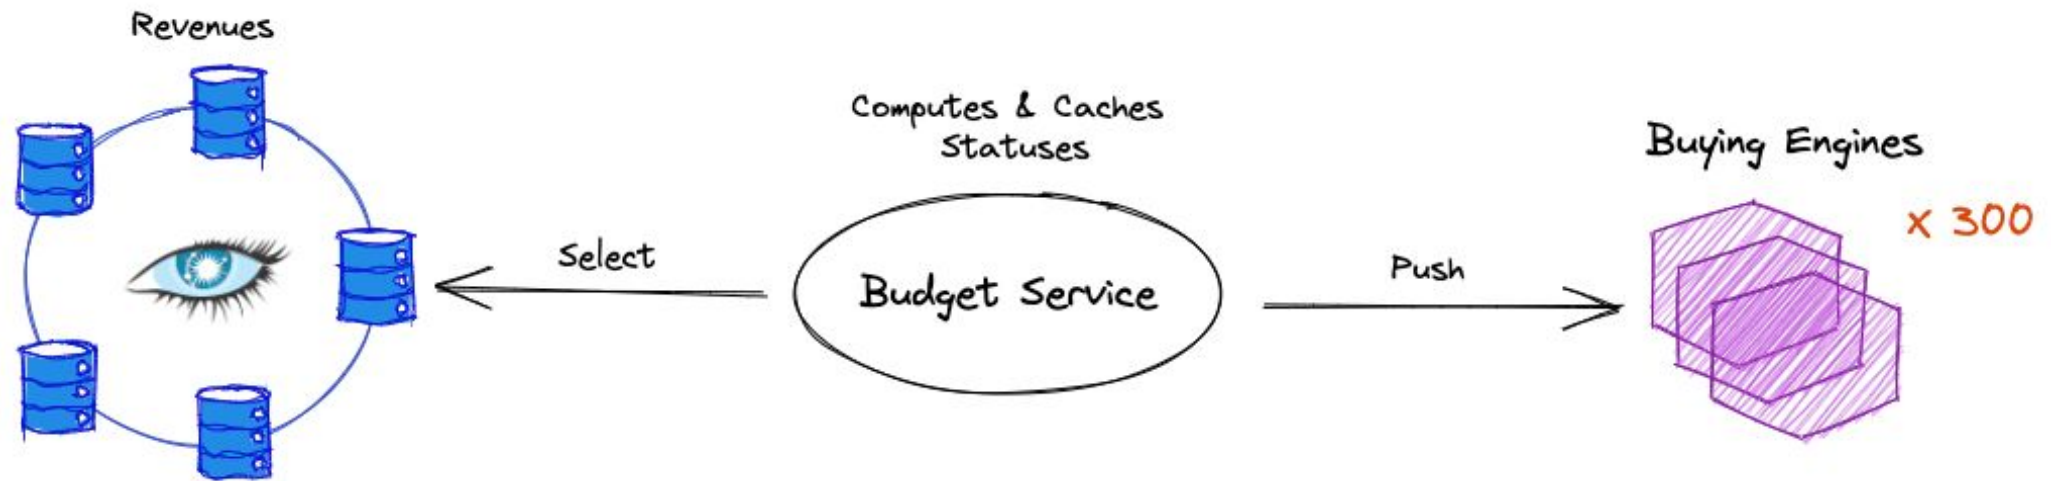
Why do we need revenue counters ?



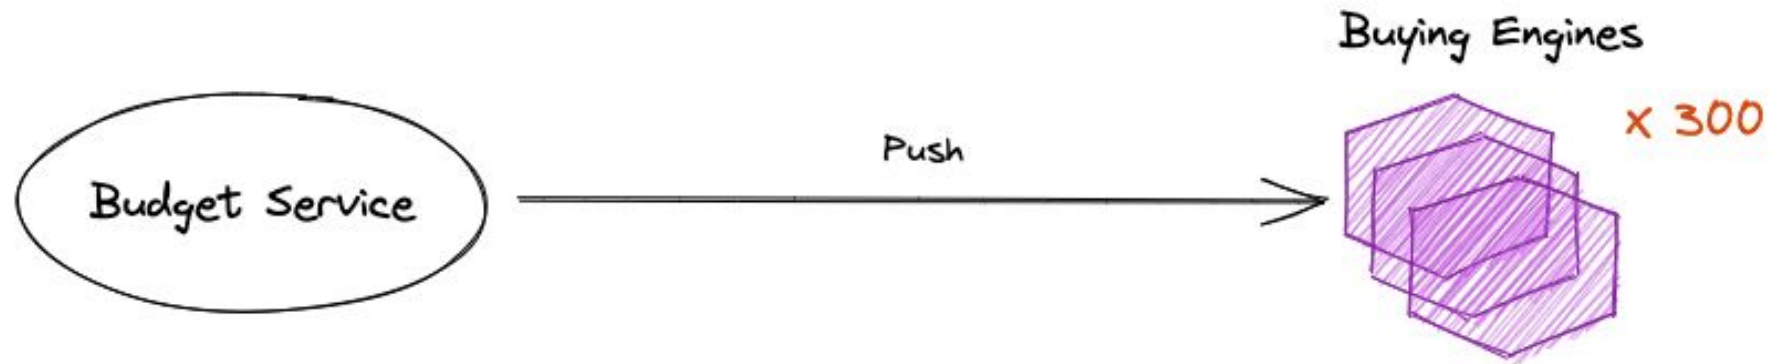
Why do we need revenue counters ?

	Data	Period	Hourly size /Ad
Revenue counters	10 doubles	2 seconds	>10000 bytes
Budget Status	1 boolean	30 minutes	~10 bytes

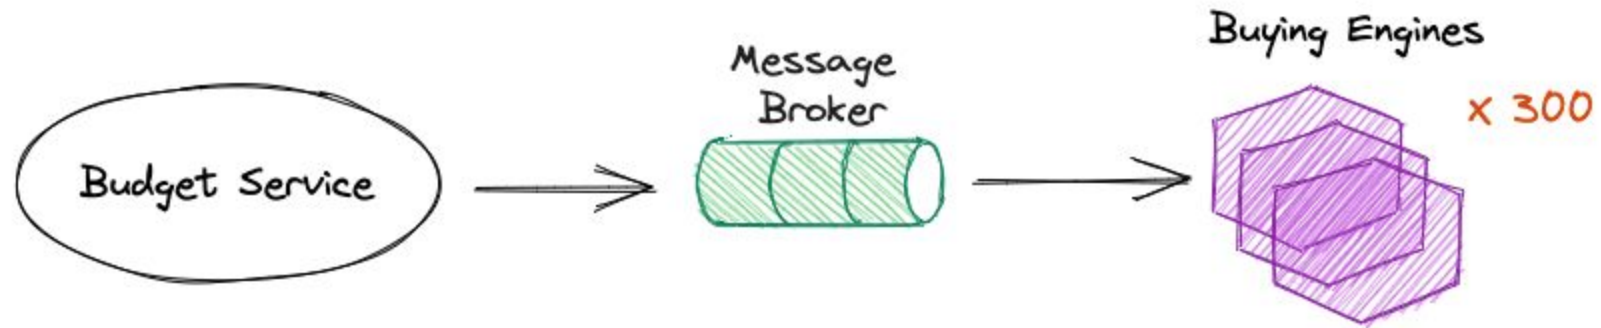
The Budget Service



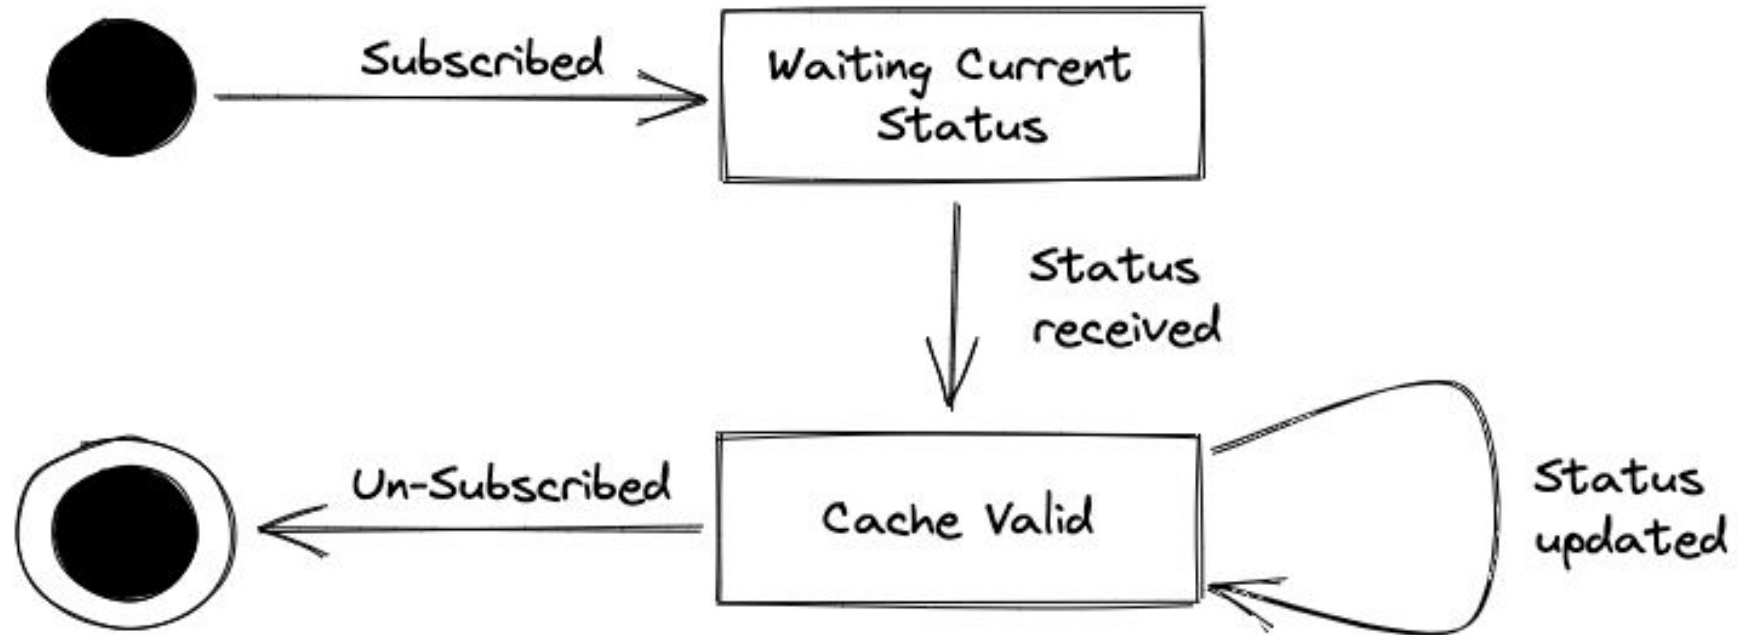
How to Push ?



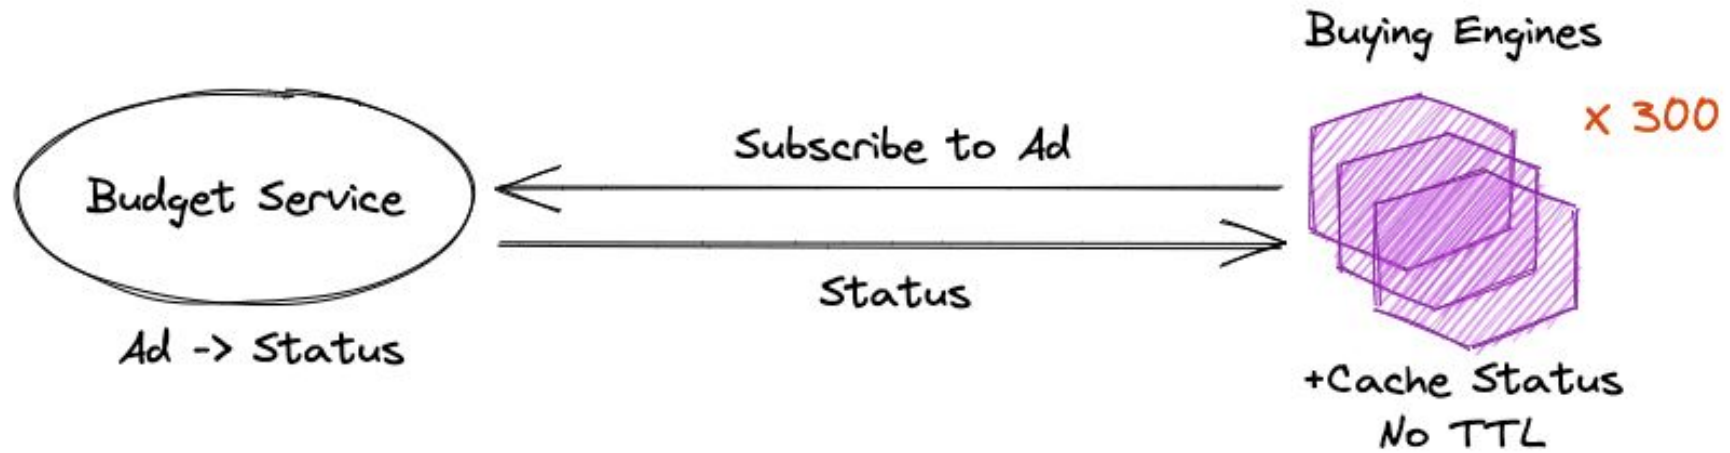
Use a Message Broker?



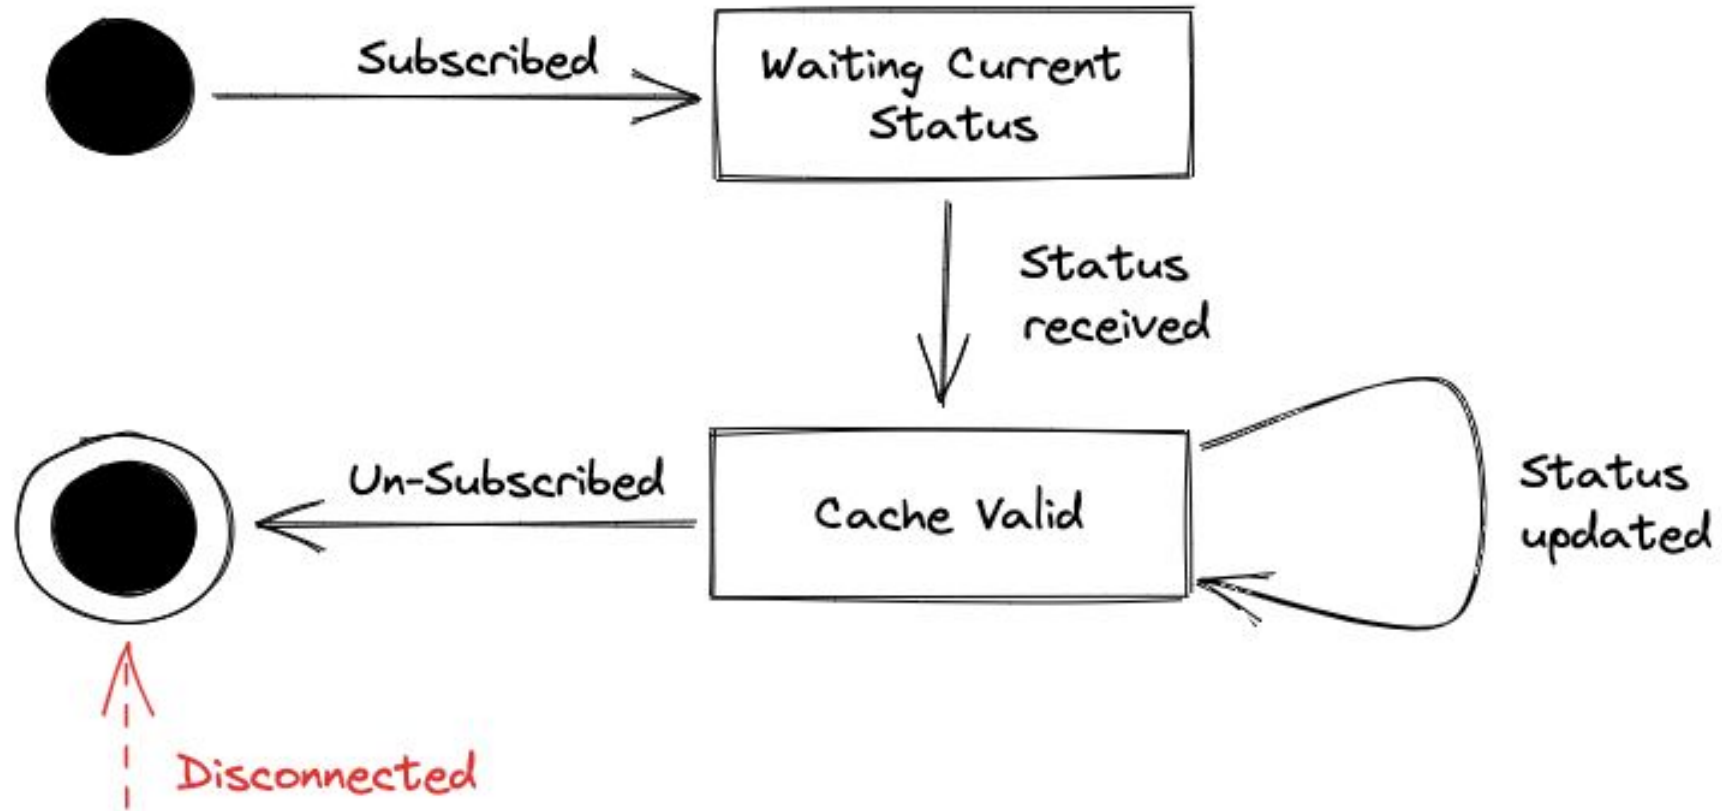
The Observer Pattern



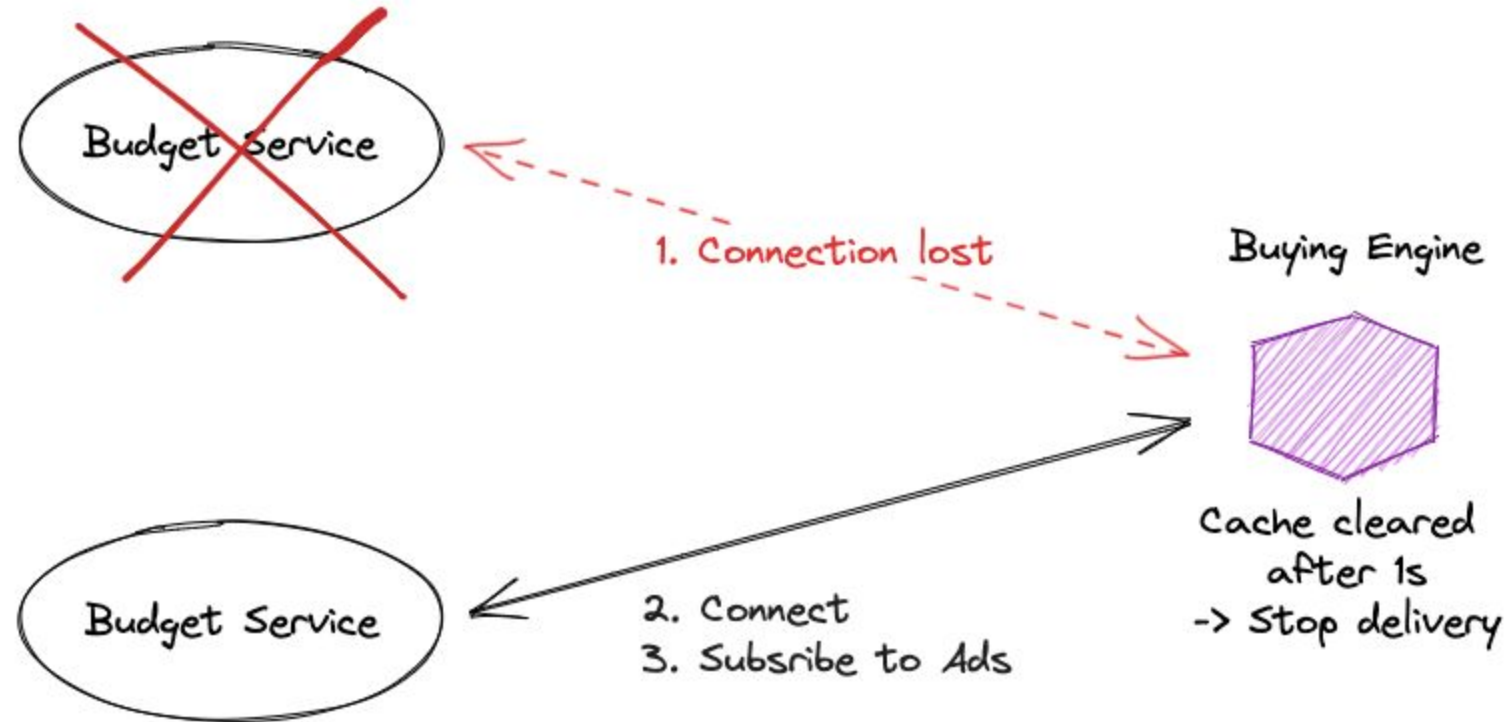
The Observer Pattern



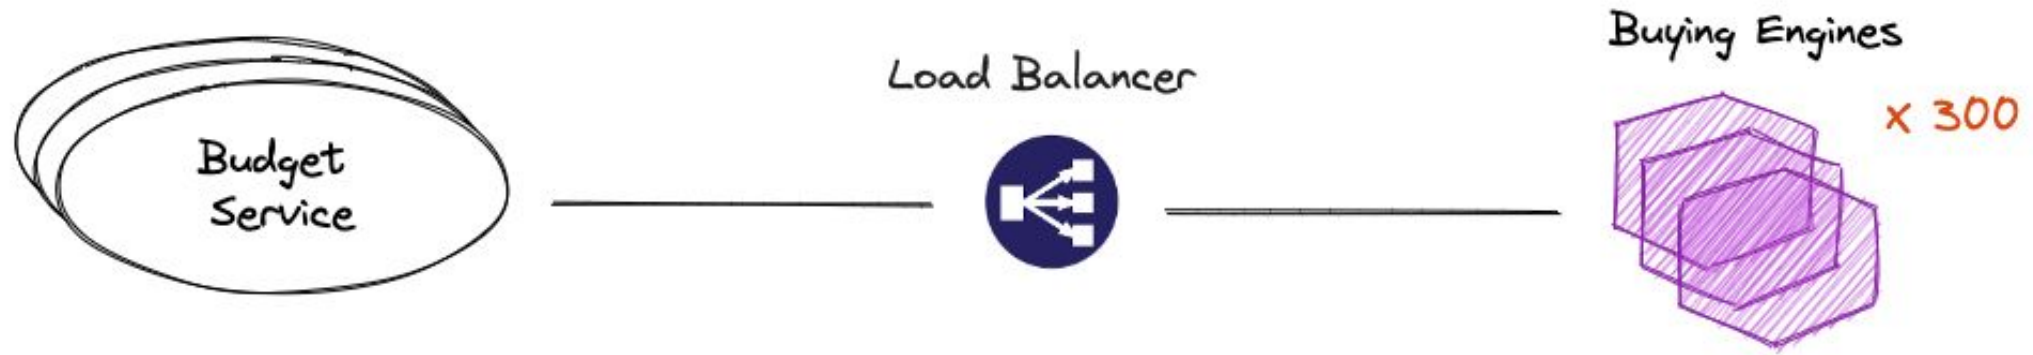
Disconnections



Disconnections



Load balancing



Implementation Choices

gRPC



- Standard
- Well supported clients
- Supports bi-directionnel streaming natively

```
service MyService {  
    rpc MyMethod (stream Request) returns (stream Response);  
}
```

Two red arrows point downwards from the 'stream' keywords in the code snippet to the 'g' and 'P' in the gRPC logo above, highlighting the bi-directional streaming capability.

gRPC

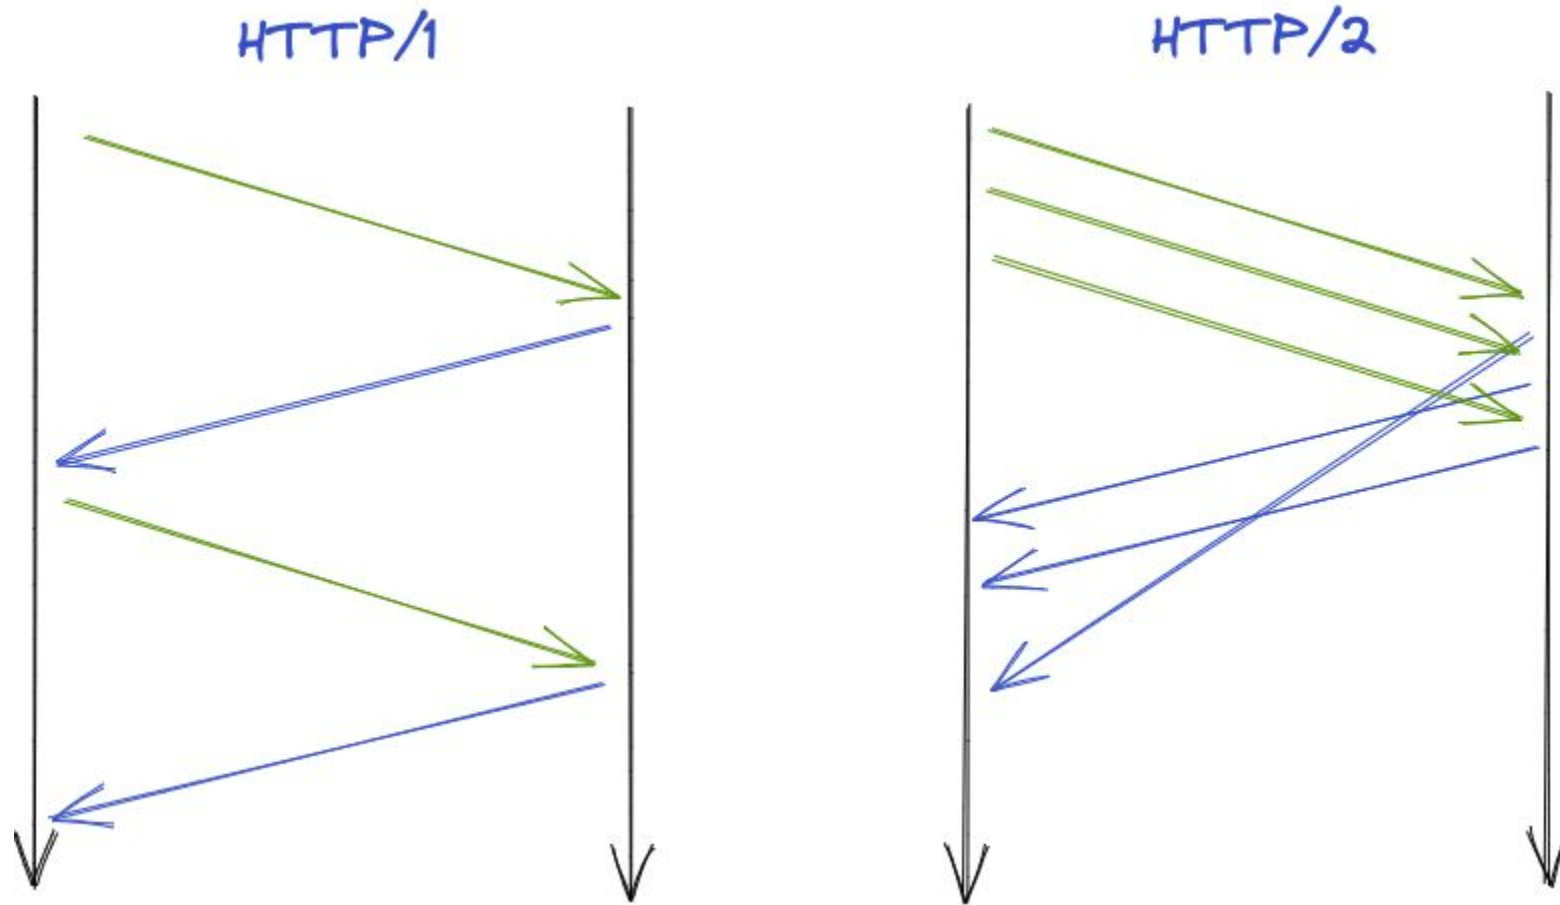


gRPC

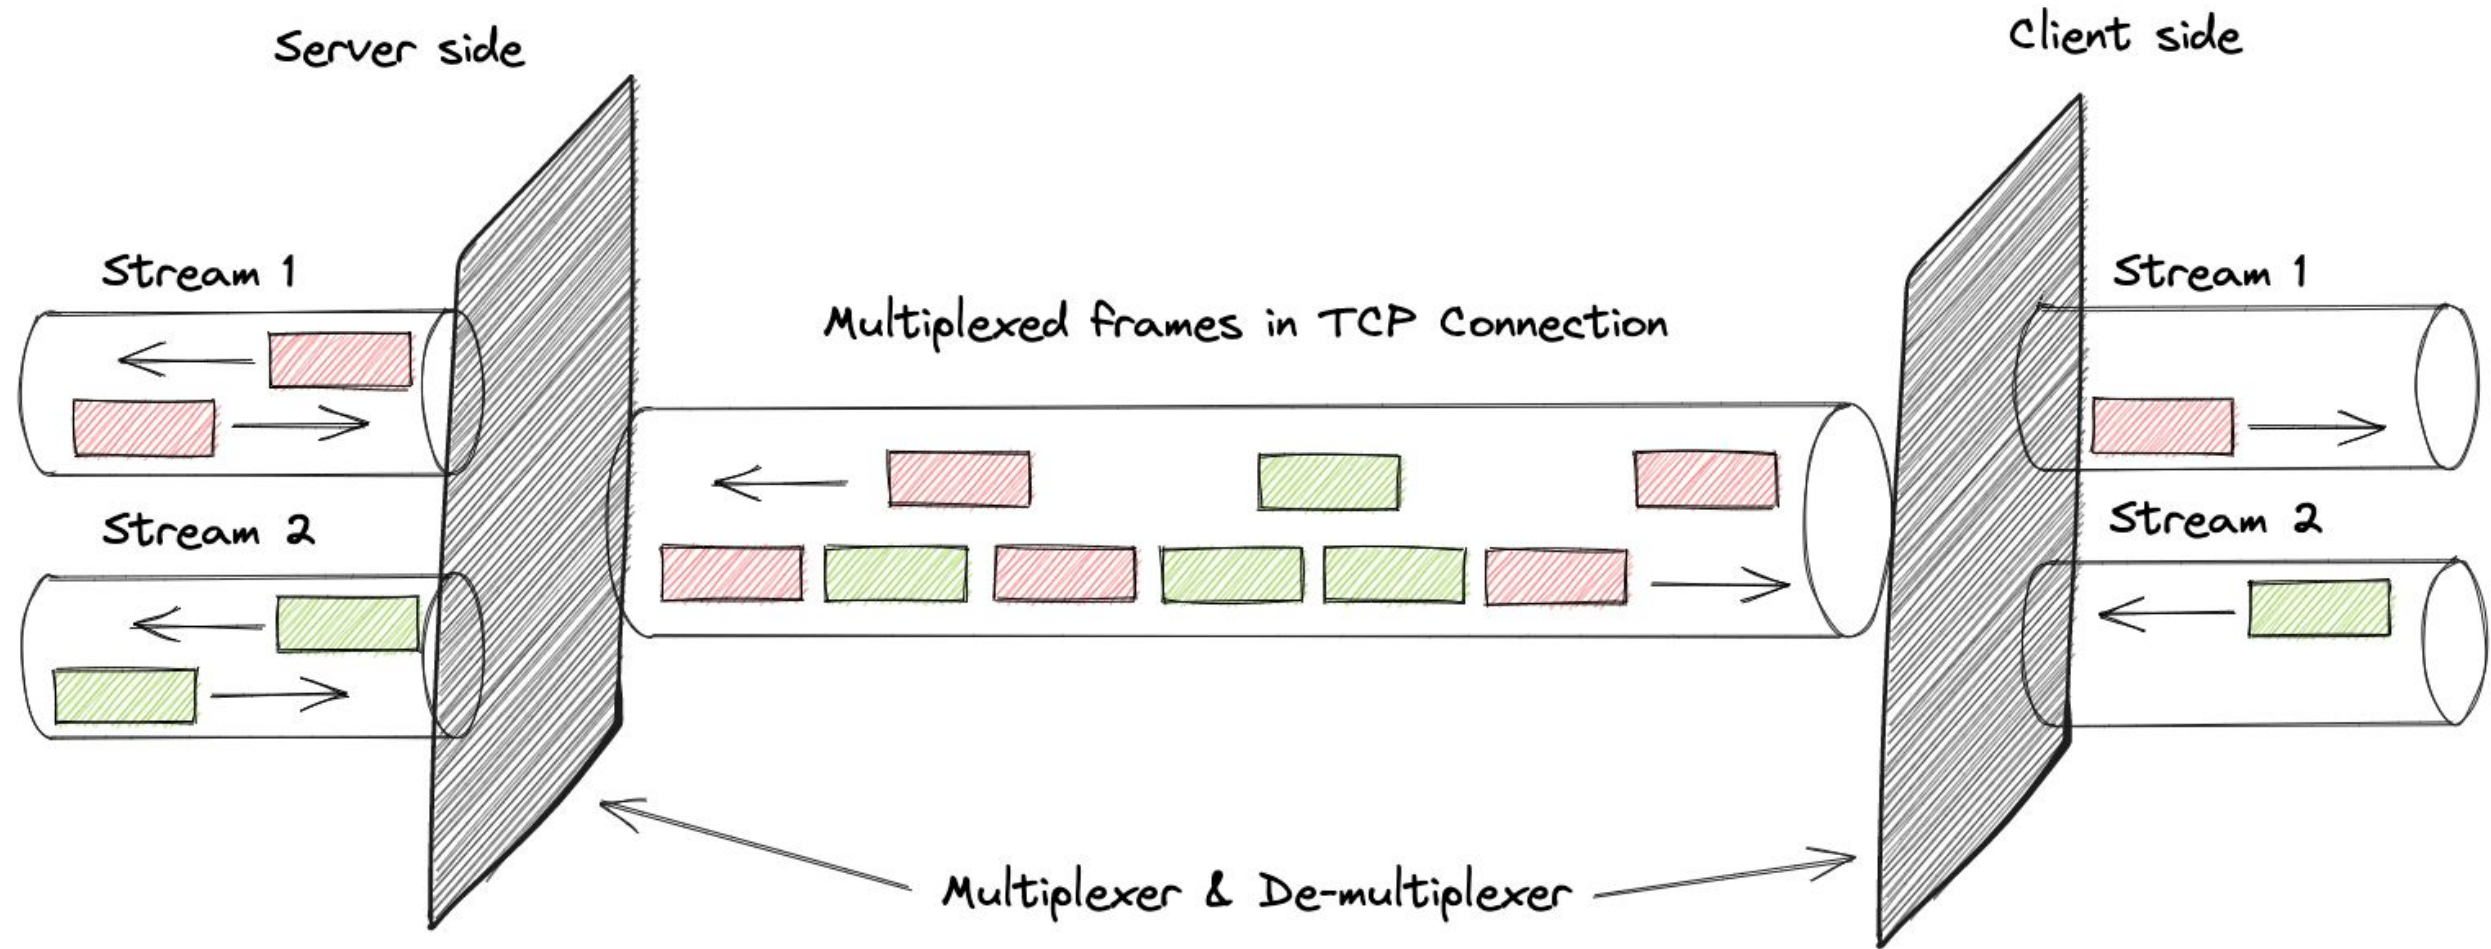
HTTP/2

TCP

Multiplexing



Multiplexing




gRPC - HTTP/2 Mapping

gRPC	HTTP/2
Channel	Connection
Call	Stream

/!\ Only one parameter

gRPC Features

- Streaming
 - Performance
 - HTTP/2
 - Protobuf
 - Eco-system (TLS, HTTP/2 servers, clients, proxies)
- 
- A decorative gradient bar at the bottom of the slide, transitioning from a light purple/pink on the left to a light blue on the right.

Service Definition

```
service BudgetService {  
    rpc SubscribeToStatus (stream SubscriptionRequest)  
        returns (stream AdStatus);  
}
```

Service Definition

```
service BudgetService {  
    rpc SubscribeToStatus (stream SubscriptionRequest)  
        returns (stream AdStatus);  
}  
  
message SubscriptionRequest {  
    int64 ad_id = 1;  
    bool subscription_toggle = 2;  
}
```


Service Definition

```
service BudgetService {  
    rpc SubscribeToStatus (stream SubscriptionRequest)  
        returns (stream AdStatus);  
}  
  
message SubscriptionRequest {  
    int64 ad_id = 1;  
    bool subscription_toggle = 2;  
}  
  
enum Status {  
    BudgetAvailable = 0;  
    BudgetExhausted = 1;  
}  
  
message AdStatus {  
    int64 ad_id = 1;  
    Status status = 2;  
}
```

Scala gRPC Implementation

gRPC Impl

[Akka-gRPC](#)

[FS2-gRPC](#)

[ZIO-gRPC](#)

Stream Library

Akka Stream

FS2

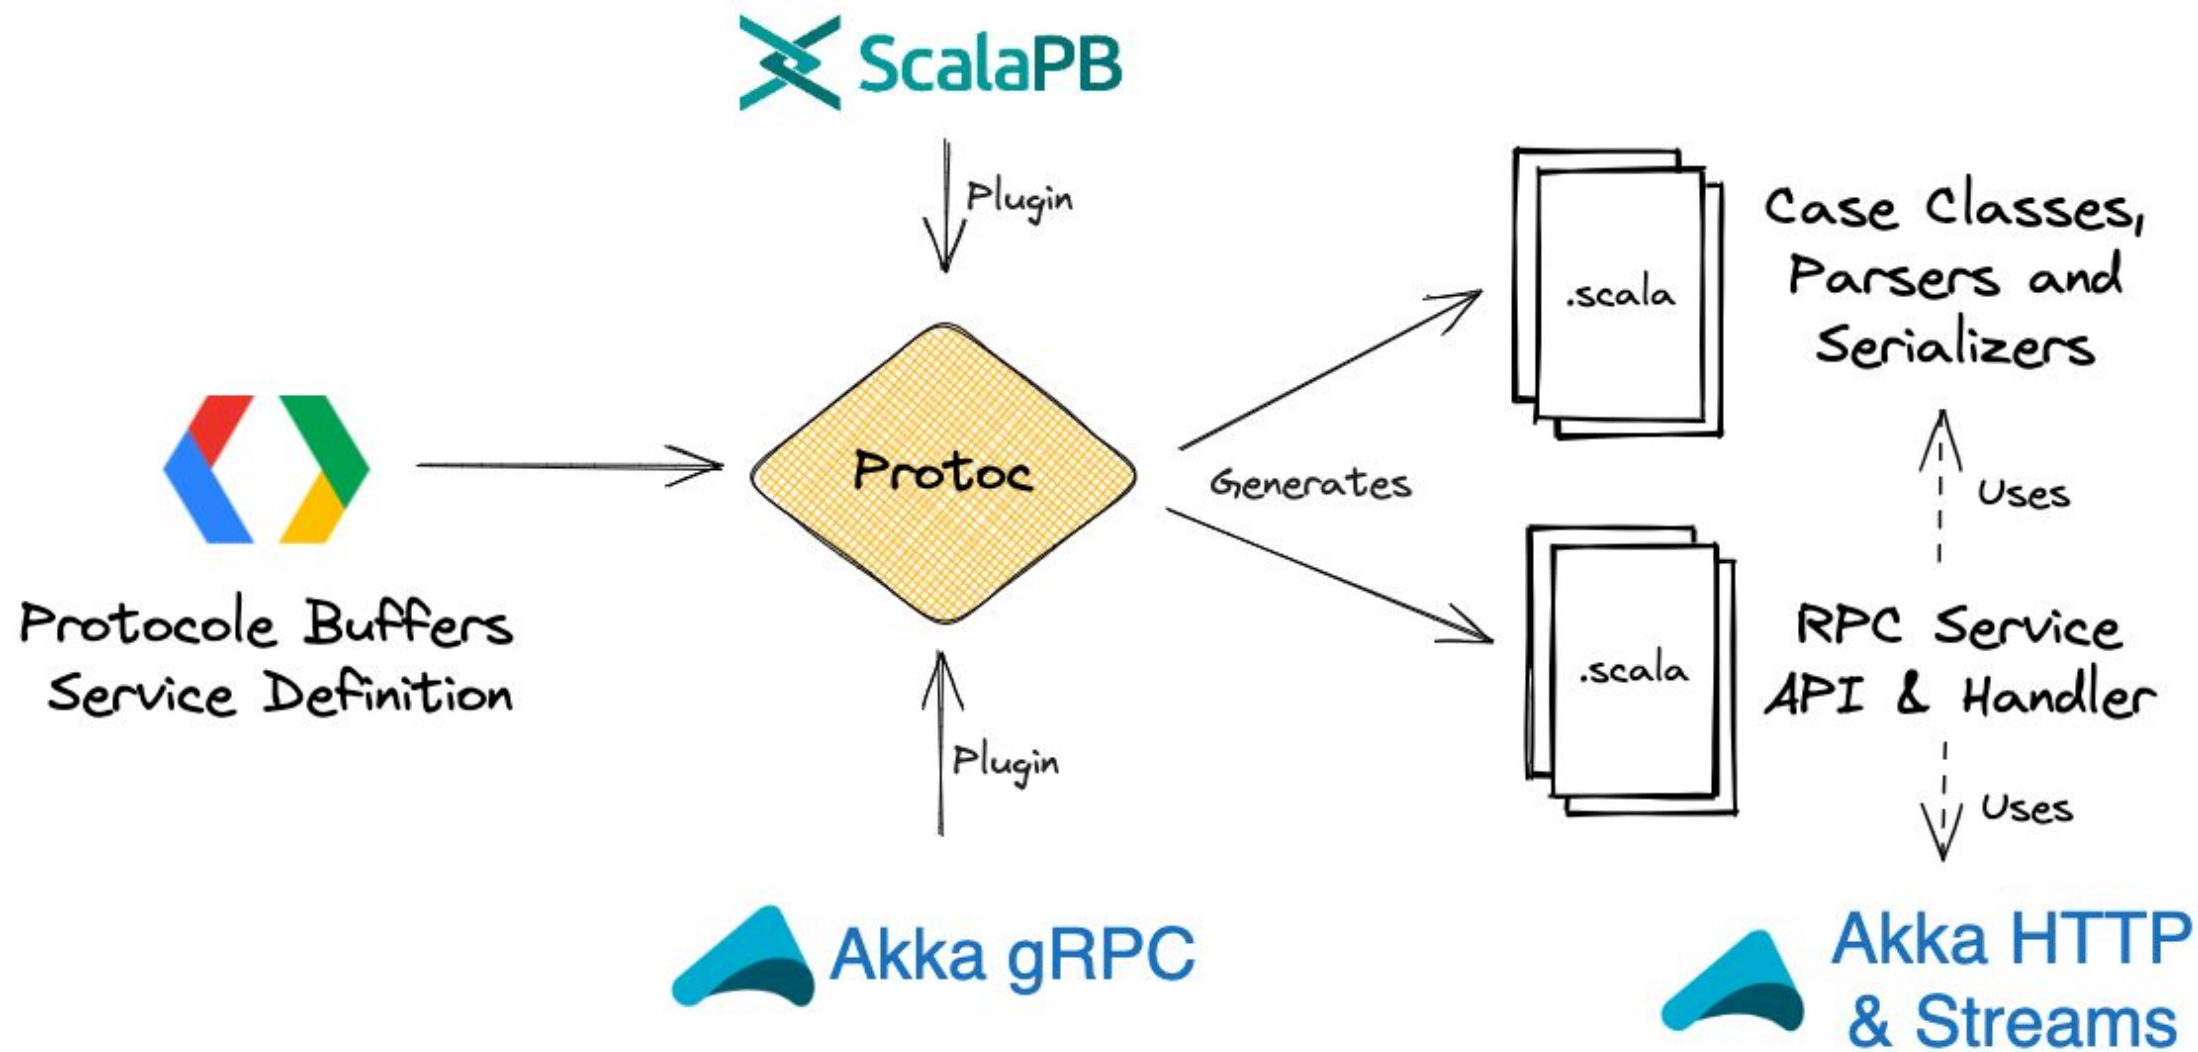
ZStream

Runtime

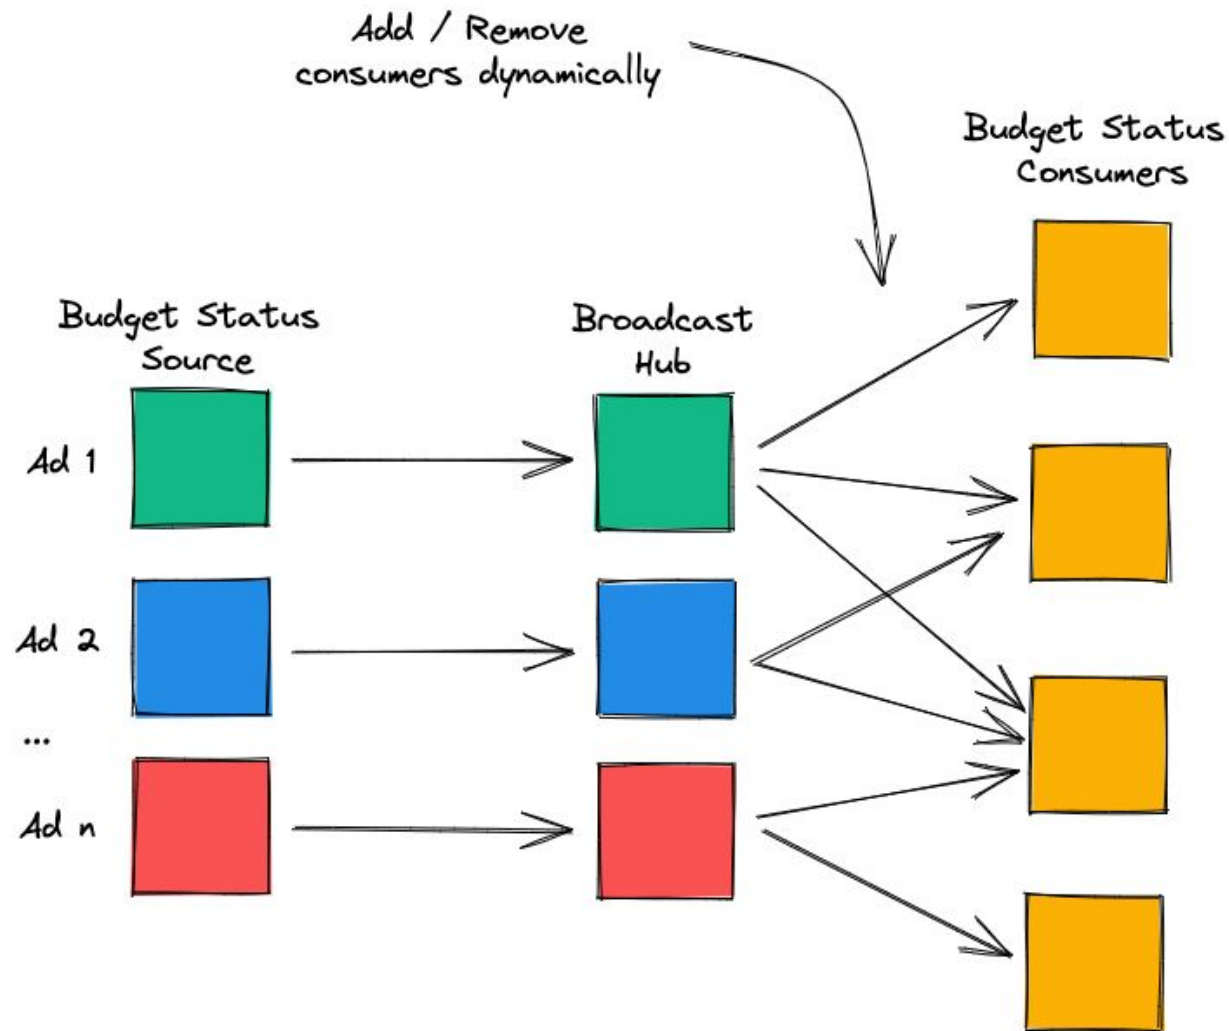
Akka

Cats Effects

ZIO



Use Akka Stream & Broadcast Hub ?



Thread Races

Ad 1 Status Computation thread

Recompute Status

Write New Status: Exhausted -
Read subscriptions -

Publish status Exhausted to each -

gRPC thread

Subscribe connection A to Ad 1

- Read current status Available

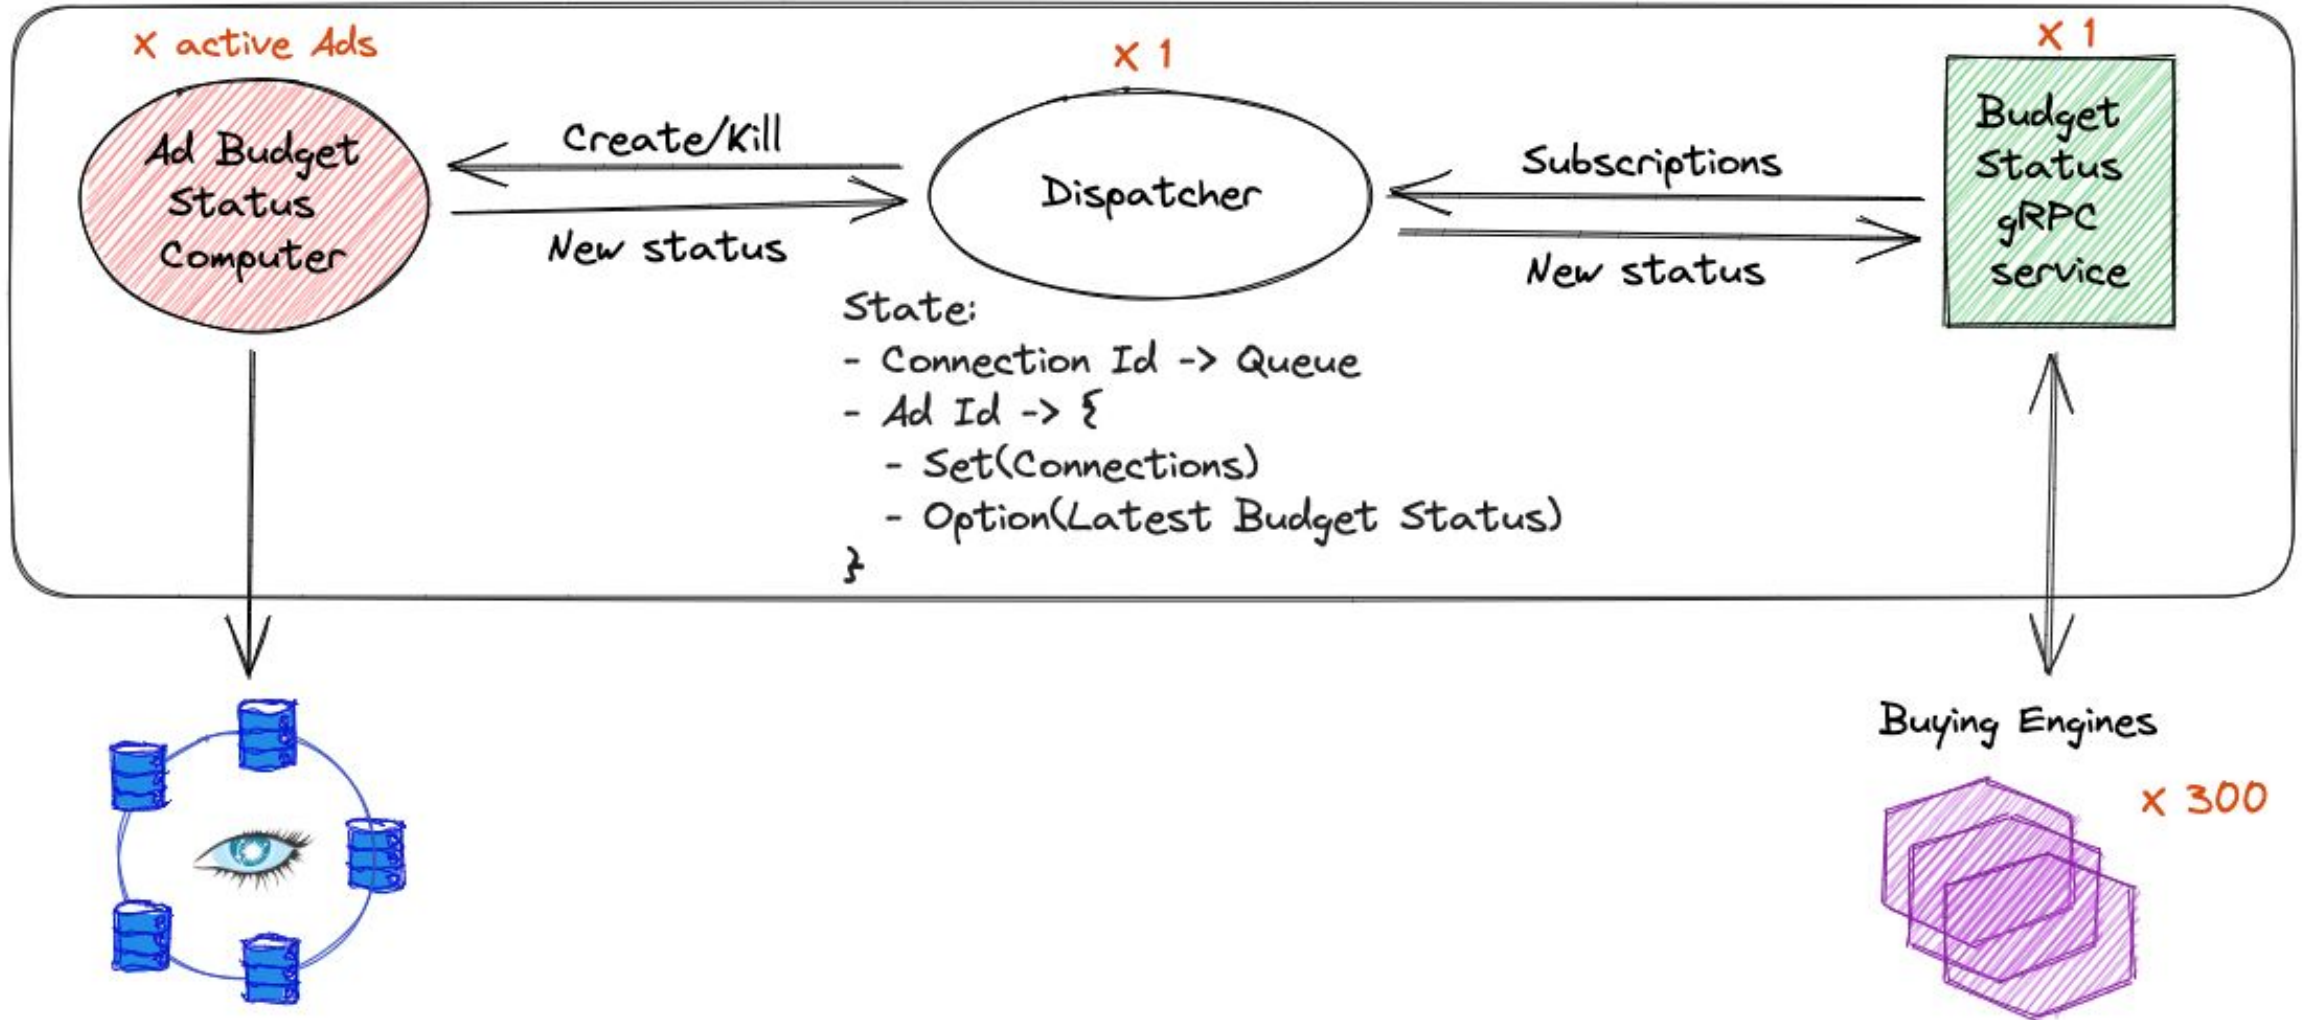
- Publish latest status Available

- Add subscription



Use Actors

Budget Service



gRPC Service Implementation

```
class BudgetServiceImpl(dispatcherRef: ActorRef[DispatcherActor.Command])  
    (implicit mat: Materializer) extends BudgetService {  
  
    private implicit val ec: ExecutionContextExecutor = mat.executionContext  
  
    override def subscribeToStatus(  
        inSrc: Source[SubscriptionRequest, NotUsed]  
    ): Source[AdStatus, NotUsed] = ???  
}
```

gRPC Service Implementation

```
val (outBoundQueue, outSrc) = Source.queue[AdStatus](1000).preMaterialize()
```


gRPC Service Implementation

```
val (outBoundQueue, outSrc) = Source.queue[AdStatus](1000).preMaterialize()  
  
dispatcherRef ! DispatcherActor.Connect(id, outBoundQueue)
```

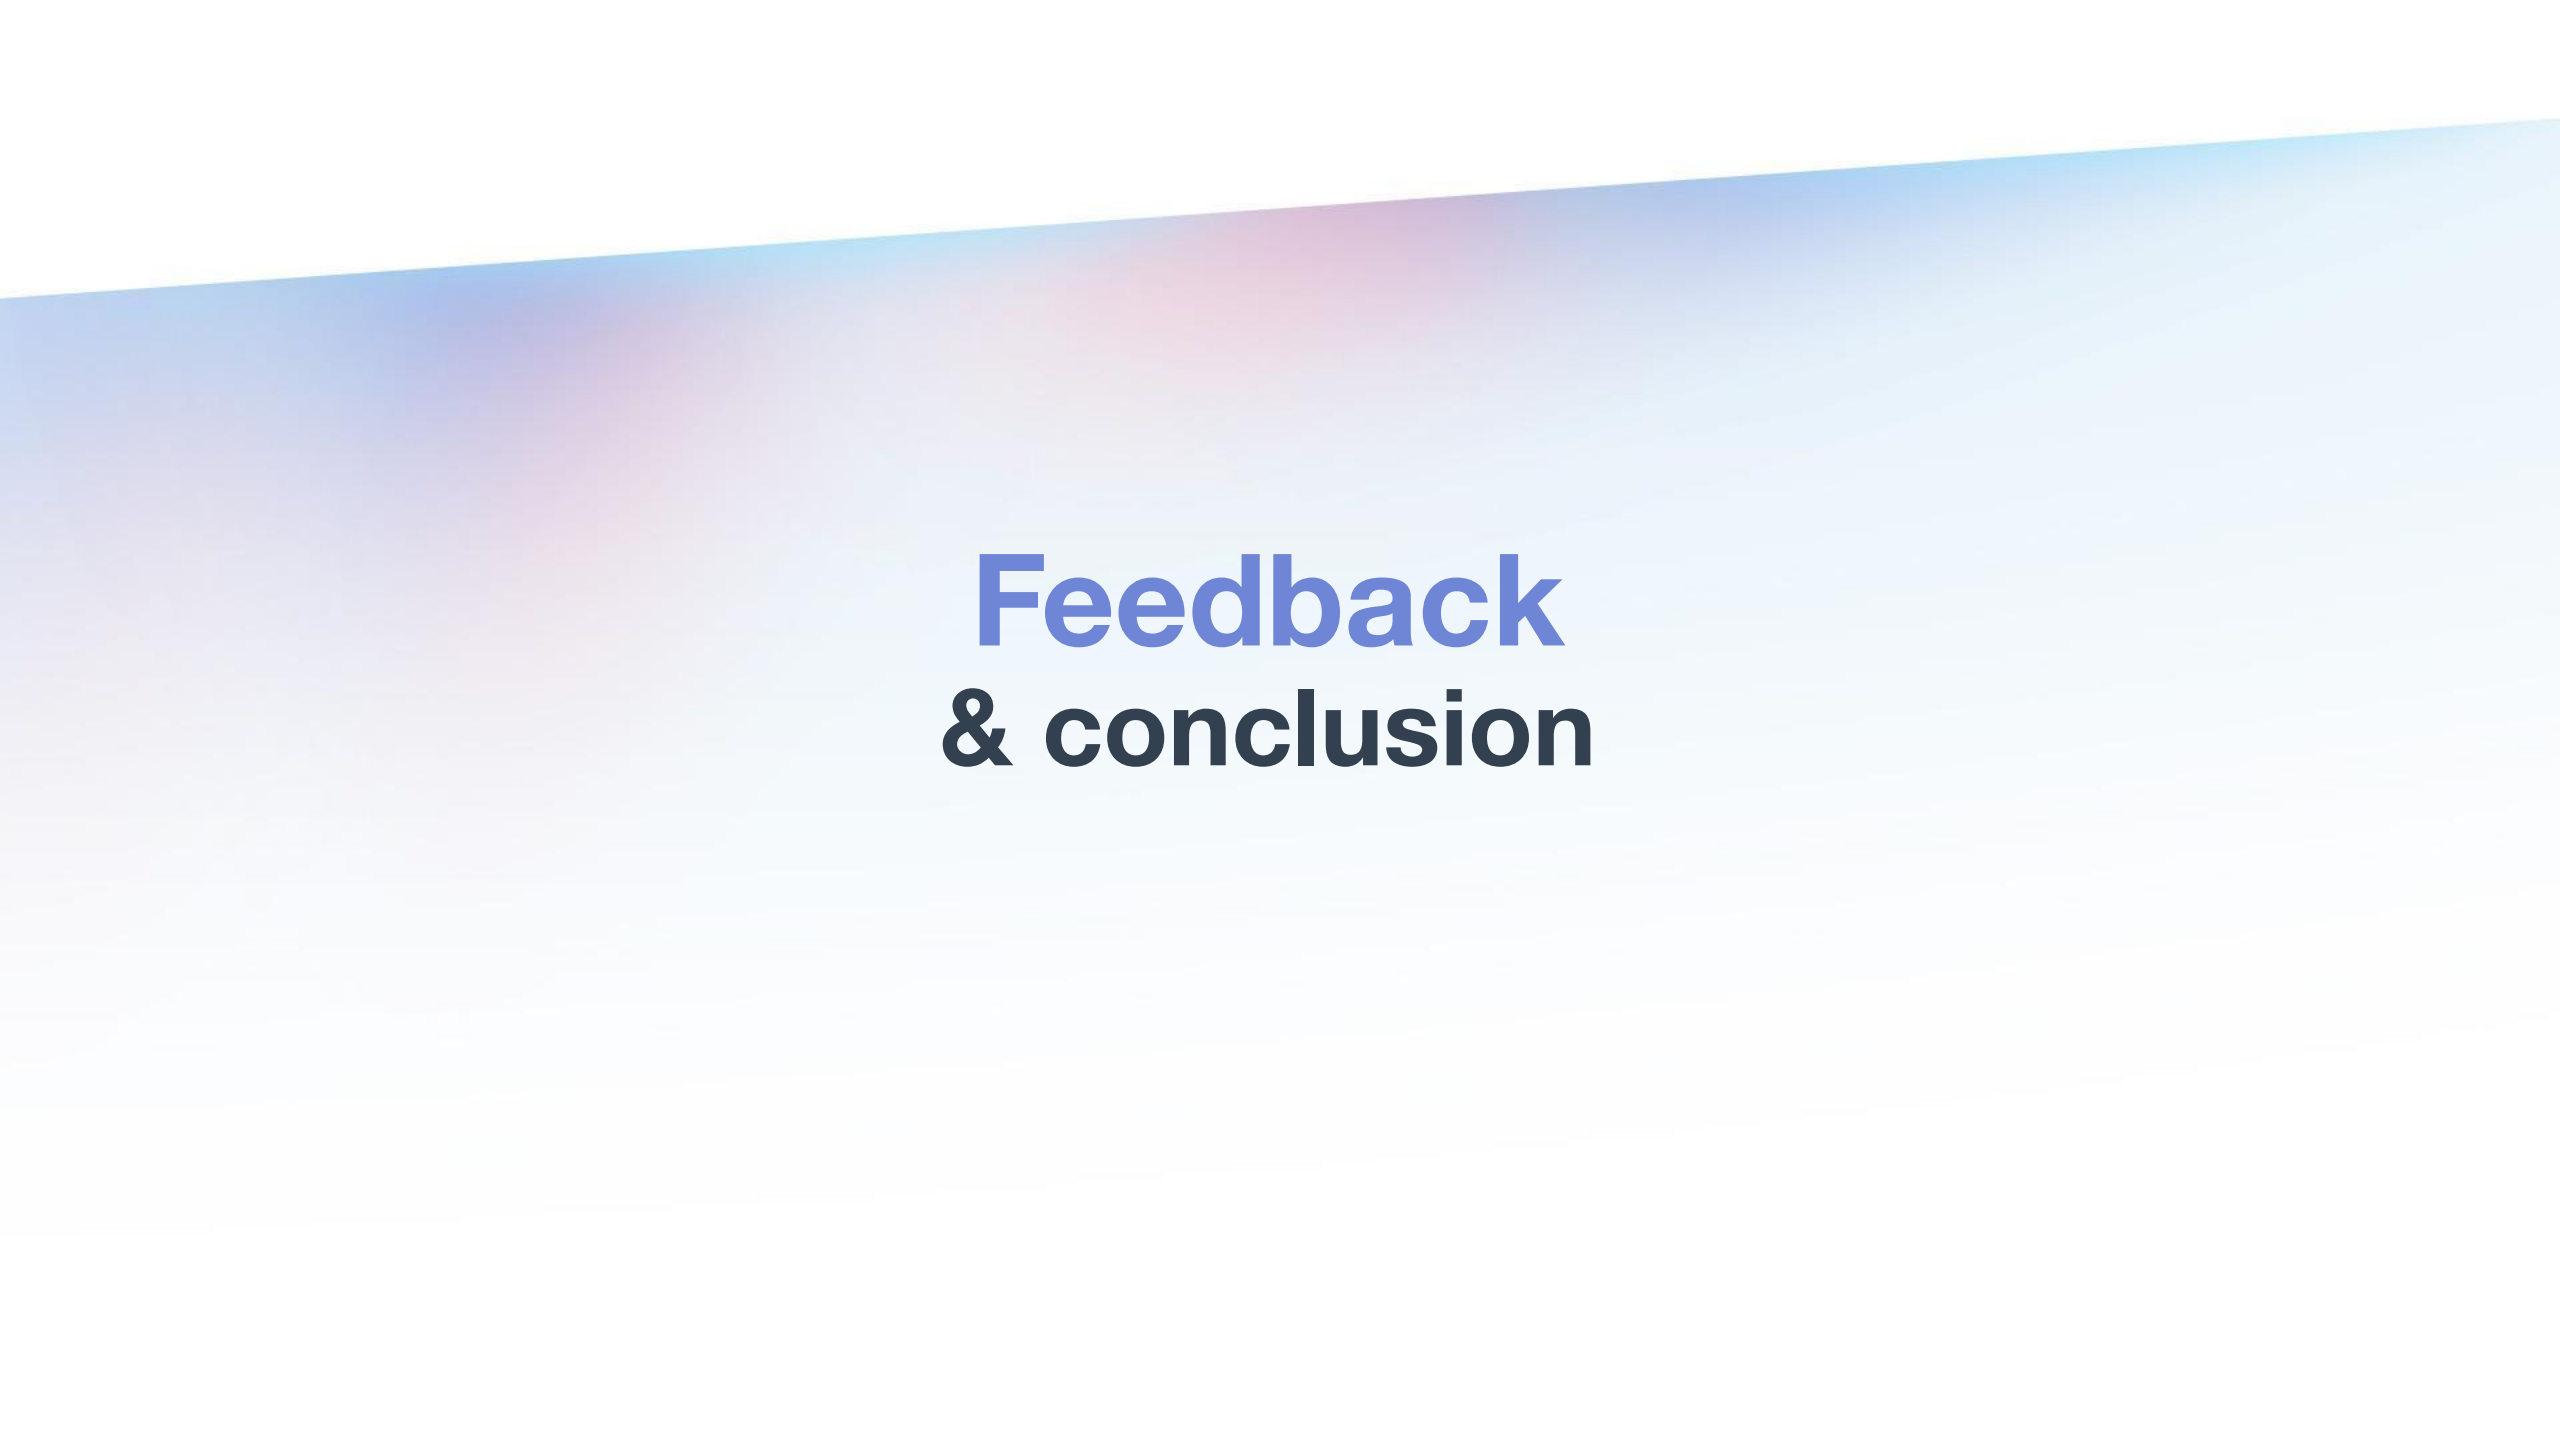
gRPC Service Implementation

```
val (outBoundQueue, outSrc) = Source.queue[AdStatus](1000).preMaterialize()

dispatcherRef ! DispatcherActor.Connect(id, outBoundQueue)


inSrc
  .runForeach { request =>
    dispatcherRef ! DispatcherActor.Subscribe(id, AdId(request.adId))
  }
  .onComplete(_ => dispatcherRef ! DispatcherActor.Disconnect(id))

outSrc
```



Feedback & conclusion

Statistics

- 2 developers
 - Allocated time: ~75%
 - Total duration: ~5 months
 - Remote observer impl: 2 weeks
- 
- A decorative gradient bar at the bottom of the slide, transitioning from a light purple/pink on the left to a light blue on the right.

Test Strategy

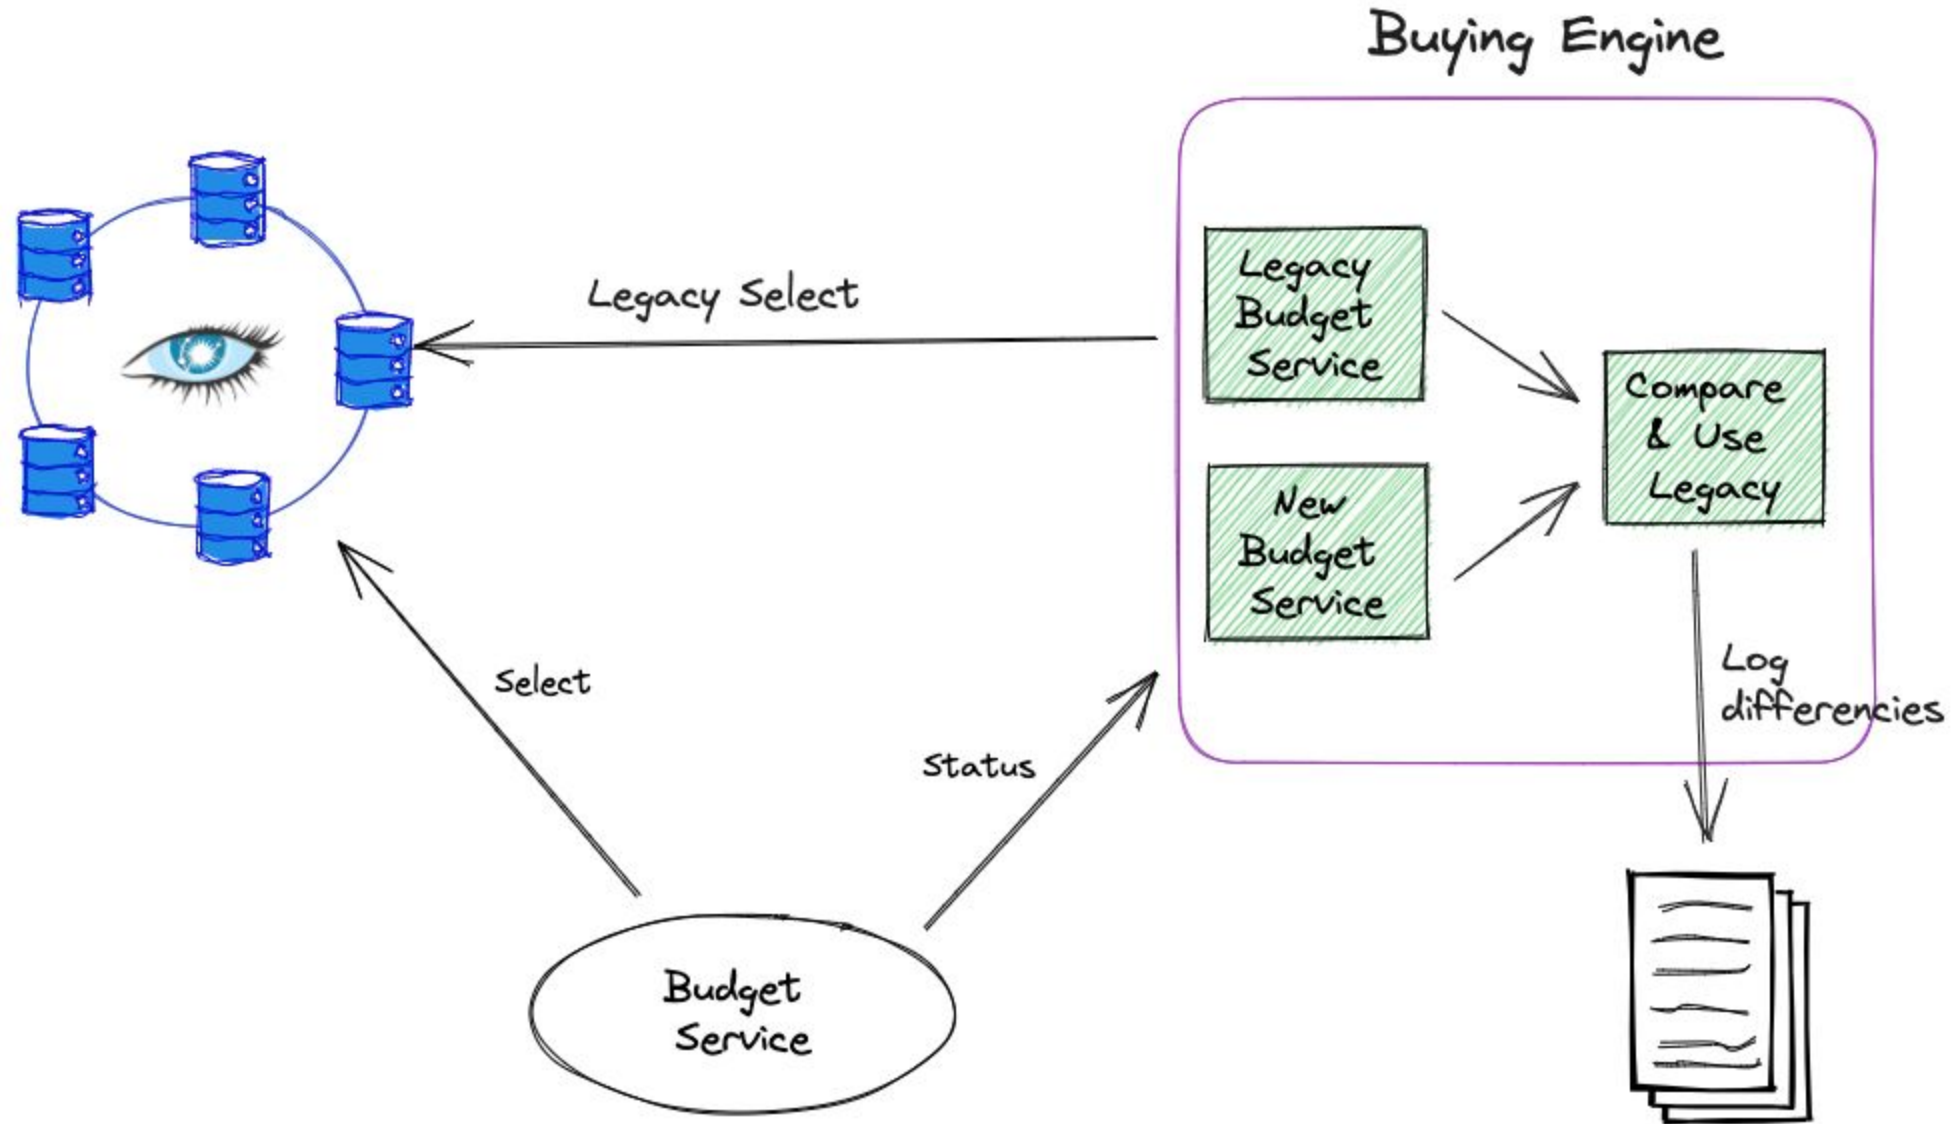
- Unitary +++
 - BDD
- Integration & Load: ---

Debug tools

- 60 metrics
 - Budget status distribution
 - 3 consecutive errors
 - Count events: subscriptions / status updates
- Alerts
 - Dispatcher lag
- /debug route

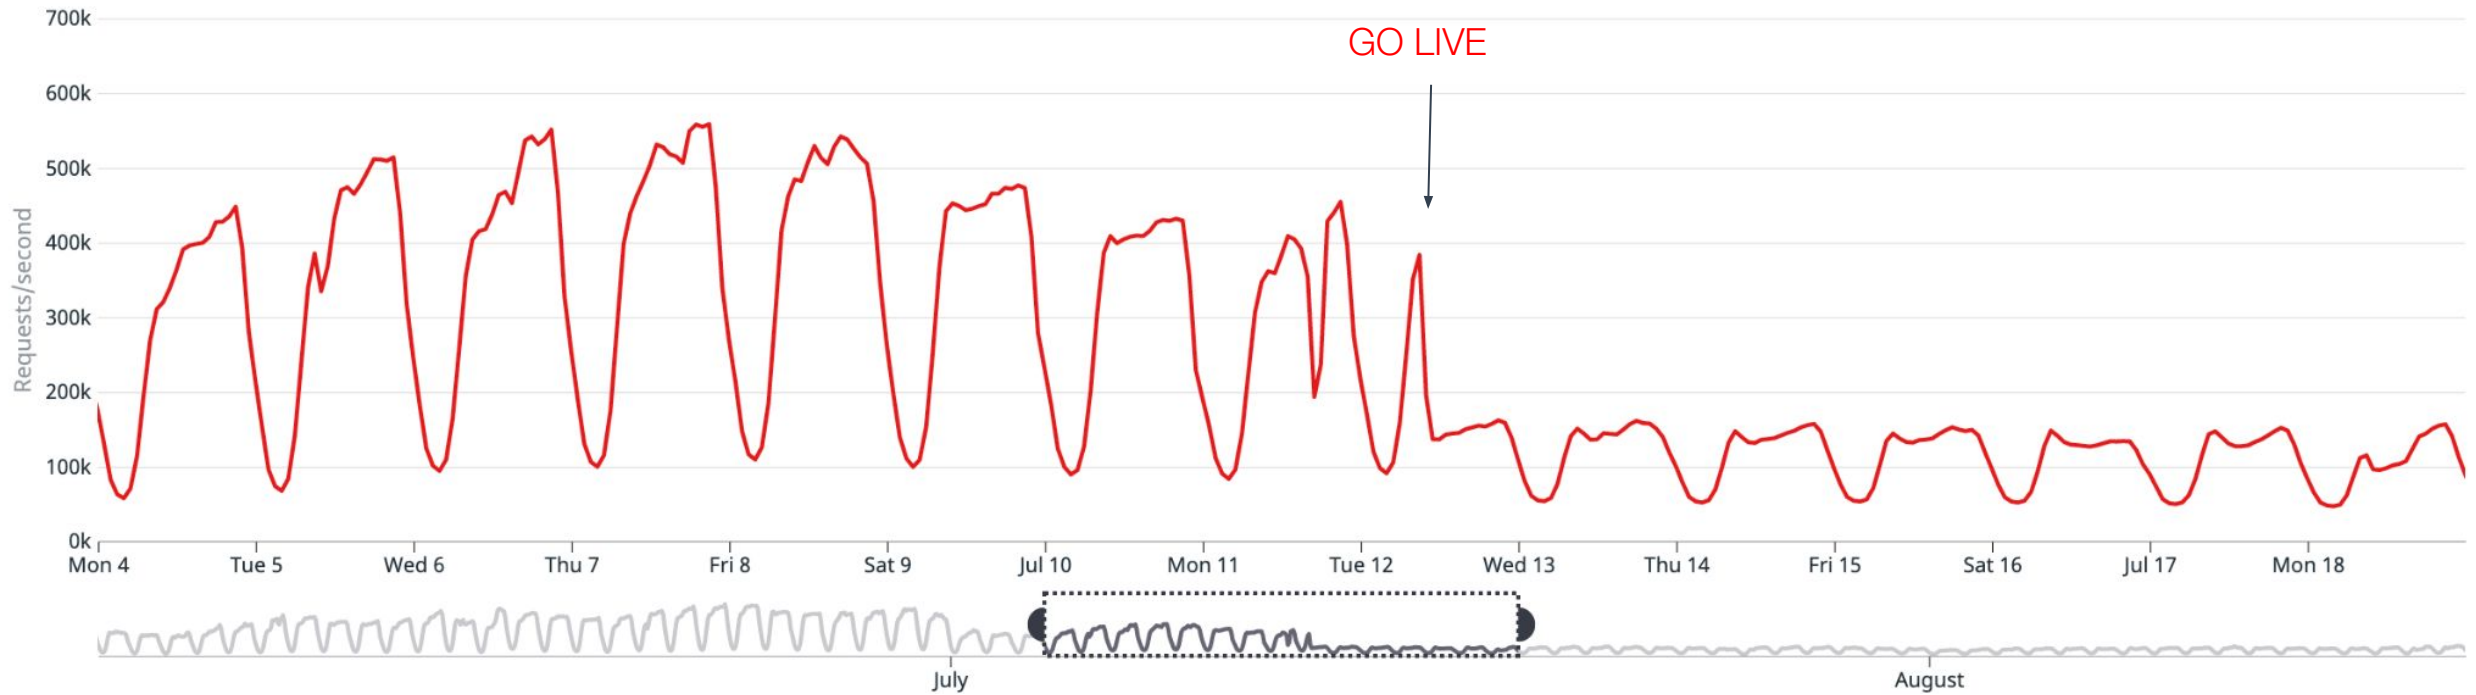


Double Run



Outcomes

Read request/s



What's next ?

- Migration vers [Pekko](#)



Takeaways

- gRPC
 - Remote Observer Pattern
 - Reliable & Sober & Simple
- 
- A decorative gradient bar at the bottom of the slide, transitioning from a light purple/pink on the left to a light blue on the right.

Useful links

- [Dynamic Cache Replication Using gRPC Streaming](#)
- [The Ad-Tech Book](#)
- [Apache Flink](#)
- [Apache Cassandra](#)
- [Google Protocol Buffers](#)
- [Scala PB](#)
- [Akka gRPC](#)
- [Apache Pekko](#)
- Illustrations were made with [excalidraw.com](#)

Questions & Responses



Feedback - - >

< - - Code sample

Illustrations made with excalidraw.com

